

Recorded Demonstration: 60-second Negative Case

Problem A claims agent | CLM-8925 | scripted backend | no API key | target length: approximately 60 seconds

Purpose

Show one negative case live: the input, the policy evidence, the escalation trigger, the actual decision log, and the deterministic Harness PASS. The teacher brief requires a five-minute team demonstration; this is the LI_LINGHAO segment.

Before recording

Open a terminal at the repository root. Do not display API keys, .env files, or unrelated private information.

```
cd /Users/llh/PycharmProjects/C-PE6201-Group-6-A2
```

60-second recording plan

Time	Open/show	Run	Must be visible
0–6 s	Terminal	<code>rm -rf ../output/recorded_demo</code> <code>python3 ../scripts/run_d4_harness.py --output-dir ../output/recorded_demo</code>	Offline scripted harness starts and completes; no key prompt.
6–14 s	Terminal	<code>SUITE=\$(find ../recorded_demo -maxdepth 1 -type d -name "suite-*" sort tail -1); echo "\$SUITE"</code>	New suite directory.
14–24 s	<code>jq ".[] select(.claim_id=="CLM-8925")" ../data_A/claims.json</code>	Three line amounts total 11,400.	
24–32 s	<code>jq ".[] select(.case_id=="CLM-8925")" ../expected_outcomes_A.json</code>	Expected escalate, trigger annual_limit_exceeded, no line pricing.	
32–46 s	<code>jq "{status,decision,trigger,escalate_to,evidence,turns,tokens_in,tokens_out,cost_usd,action_count,limits}" "\$SUITE/CLM-8925-trial-1.json"</code>	Actual escalation, correct trigger, one gated write, step cap 8.	
46–54 s	<code>jq ".trace[] {turn,tool,observation,status}" "\$SUITE/CLM-8925-trial-1.json"</code>	get_claim; lookup_policy shows remaining 9,200 and exceeded; issue_decision_letter.	
54–60 s	<code>jq "{trials,cases,negative,code_passed,execution_errors,step_cap_hits,backend,prompt_version}" "\$SUITE/summary.json"</code>	75 trials; 45 negative; 75/75 code passes; 0 errors; 0 cap hits; scripted.	

Decision log (optional if time permits)

```
jq "{decision_log,action_records}" "$SUITE/CLM-8925-trial-1.json"
```

Show exactly one issue_decision_letter record: escalate; annual_limit_exceeded; human claims assessor; evidence citations observed successfully; gate operator approved.

Exact English speech (~60 seconds)

This is negative case CLM-8925. The claim has three lines totalling 11,400, while the policy has only 9,200 remaining. The policy tool returns annual_limit_status=exceeded, so the agent stops before pricing individual lines. It escalates to a human claims assessor with the single trigger annual_limit_exceeded. The trace shows successful claim retrieval, duplicate and member checks, the policy observation, and one gated decision write. The decision log matches the final output. The deterministic scripted harness then reports 75 out of 75 code checks passed, with 45 negative trials, zero execution errors, and zero step-cap hits. This is an offline reproducible harness result, not a live-model accuracy score.